

Forfattere:

13.07-89 Rane Eriksen BVL193
05.01-90 Christian Nielsen BNF287
14.06-79 Sebastian O. Jensen GJX653
Hold 1

DATALOGISK INSTITUT

KØBENHAVNS UNIVERSITET

OPERATIVSYSTEMER OG MULTIPROGRAMMERING

G-Assignment 3

28. FEBRUAR 2015



UNIVERSITY OF
COPENHAGEN

Indhold

1	Task 1: A thread-safe unbounded queue	2
	implementation	2
	Tests	2
2	Task 2: Userland semaphores for Buenos	2
	Where we made changes	2
	Tests	7

1 Task 1: A thread-safe unbounded queue

implementation

Tests

2 Task 2: Userland semaphores for Buenos

Where we made changes

In this section we will only focus on where we made changes in order to implement our solution to userland semaphores.

We will explain what we implemented and why later, this section is only to create a simple overview of what once was and now is.

init/main.c

This section initializes the userland semaphores along with every other initialization in init/main.c

```
kwrite("Initializing userland semaphores\n");  
userland_semaphore_init();
```

proc/syscall.h

We added this line because in order to follow the standards of the current buenos

```
#define SYSCALL_SEM_CLOSE 0x303
```

proc/syscall.c

Add something to this section!

tests/lib.h

In order to avoid name clashes of function in different areas of the buenos we renamed a few functions in tests/lib.c. The changes are as follows:

‘putc’ was changed to ‘putc_userland’
‘strlen’ was changed to ‘strlen_userland’

The functions below are included in this header file and are explained in detail in the next section

```
usr_sem_t* syscall_sem_open(char const* name, int value);
```

```
int syscall_sem_p(usr_sem_t* handle);
```

```
int syscall_sem_v(usr_sem_t* handle);
```

```
int syscall_sem_destroy(usr_sem_t* handle);
```

tests/lib.c

In order to avoid name clashes of function in different areas of the buenos we renamed a few functions in tests/lib.c. The changes are as follows:

‘putc’ was changed to ‘putc_userland’

‘strlen’ was changed to ‘strlen_userland’

In this file we also added four functions that are visible to any program programmed to run in the buenos OS and acts as a link between the operating system and user program.

The four functions are as follows:

A system call to this function enables a user program to create a semaphore or access a semaphore that is already created. How this function works is explained in subsection proc/semaphore.c

```
usr_sem_t* syscall_sem_open(char const* name, int value) {  
    return (usr_sem_t*)_syscall(SYS_CALL_SEM_OPEN,  
                                (uint32_t)name, (uint32_t)value, 0);  
}
```

A system call to this function enables a user program or process to block itself untill another program or process releases the block. How this function works is explained in subsection proc/semaphore.c

```
int syscall_sem_p(usr_sem_t* handle) {  
    return (int)_syscall(SYS_CALL_SEM_PROCURE,  
                        (uint32_t)handle, 0, 0);  
}
```

A system call to this function enables a user program or process to release another program or process that has blocked itself by the calling syscall_sem_p. How this function works is explained in subsection proc/semaphore.c

```
int syscall_sem_v(usr_sem_t* handle) {  
    return (int)_syscall(SYSCALL_SEM_VACATE,  
                        (uint32_t)handle, 0, 0);  
}
```

A system call to this function enables a user program or process to close a semaphore if no other program or process is blocked by it. How this function works is explained in subsection `proc/semaphore.c`

```
int syscall_sem_destroy(usr_sem_t* handle) {  
    return (int)_syscall(SYSCALL_SEM_CLOSE,  
                        (uint32_t)handle, 0, 0);  
}
```

proc/semaphore.h

We took what was already implemented in `kernel/semaphore.h` and modified the existing code to serve the needs to hold a name as identifier for the semaphore and to bring a semaphore to buenos userland which is not related to the semaphores used by buenos kernel.

Instead of having creator point to a thread we made it point to the process which created the semaphore - the creator id is strictly used to see if the semaphore is in use or if it's available.

The maximum length of the semaphore is currentl 256, but this is only due to execute or tests that needs the name to be minimum 129 characters long. The maximum length of the name of the semaphore should not need to be longer than 128 or 64 characters long because we estimate that developers would rather keep short and precise names rather than long names in the programming code.

```
#ifndef BUENOS_USERLAND_SEMAPHORE_H  
#define BUENOS_USERLAND_SEMAPHORE_H  
  
#define MAX_SEMAPHORES 128  
#define MAX_SEMAPHORE_NAME 256  
  
#define FAILURE -1  
  
#include "kernel/spinlock.h"  
#include "proc/process.h"
```

```
typedef struct {
    char name[MAX_SEMAPHORE_NAME];
    spinlock_t slock;
    int value;
    process_id_t creator;
} usr_sem_t;

void userland_semaphore_init(void);
usr_sem_t *userland_semaphore_create(char const* name, int value);
int userland_semaphore_P(usr_sem_t *sem);
int userland_semaphore_V(usr_sem_t *sem);
int userland_semaphore_destroy(usr_sem_t *sem);

#endif /* BUENOS_USERLAND_SEMAPHORE_H */
```

proc/semaphore.c

In this section we will only include the most essential parts that differ from the kernel semaphore in an attempt to keep this report short and precise.

As stated in the previous subsection we took the semaphore from the kernel and modified it to fulfill the need to hold a name.

The initializer function does the exact same thing as the kernel version except that it sets the first character of the name to the null terminating string.

‘userland_semaphore_create’ is split up in two functions and the name may be misleading as name, but it is divided in to two functions: ‘find_existing_semaphore’ and ‘create_fresh_semaphore’ are the two functions that actually defines ‘userland_semaphore_create’, the decision between the two functions are based on the value that is passed to ‘userland_semaphore_create’. If the value is negative ‘find_existing_semaphore’ is called and if the value is equal to or higher than zero ‘create_fresh_semaphore’ is called.

‘create_fresh_semaphore’ does as titled, creates a new semaphore if one is available however it needs to check every entry in the table to make sure that no other semaphore by the name given already exists, the algorithm therefore runs in $\Theta(n)$ in worst case which makes it expensive, the reason it’s expensive compared to the one in the kernel which also runs $\Theta(n)$ in worst case is, that the one we have implemented must keep checking for existing semaphores with the name provided even though it has determined that a

semaphore is available.

A way to make it faster would be to index the names in alphabetically order which would drastically speed up the function however indexing might be time consuming and there's always the case where a programmer always starts the name of the semaphores with a prefix which would then affect the indexing of the names in a worse way, also there is the consideration of the cost when actually indexing the names which could also be expensive.

If the function created a semaphore the pointer to the semaphore is returned else NULL is returned.

'find_existing_semaphore' also does as titled, it searches for a semaphore with the given name the running time is also $\Theta(n)$ however it has only one condition, which makes it faster in most cases.

If the semaphore exists the pointer to the semaphore is returned otherwise NULL is returned

'userland_semaphore_P' the only difference from this function and the function used by the kernel is, that we do not trust programmers, we therefore test for an initialized semaphore. If the semaphore 'creator' value is -1 the function returns -1 to indicate an error otherwise 0.

There's a possibility that the function simply does not exist in memory, however TLB exception should handle this case.

'userland_semaphore_V' differs from the kernel version in the same way the 'find_existing_semaphore' differs from kernel.

The return value is the exact same as well under the same circumstances.

Lastly comes 'userland_semaphore_destroy' which enables a program to close a semaphore thus making it available to other programs, it is very important that a program or process closes every semaphore created in order to prevent a memoryleak since buenos will not destroy these as the process is terminated in any way.

A semaphore can be used across different processes which is why we did not implement termination for every semaphore once its process terminates - This is left to the programmer.

Tests

In order to test our implementation we have created a rather large test program consisting of three programs:

`semaphore_test.c`
`semaphore_test2.c` and
`semaphore_test3.c`.

Before we discuss our tests we want to make clear what we want to test.

Regarding creating or reopening a semaphore we want to test that we get the correct return values based on the input we give - We want to test that when we call `'userland_semaphore_create'` with a given name and a value that is at least 0 or higher that we are given a new semaphore and in case of errors that we get NULL.

The same test goes for when the value is negative.

We also want to test that when `'userland_semaphore_P'` the program or process calling the function is blocked, if the value is less than 0.

Next to this we want to test that calling `'userland_semaphore_V'` results in unblocking one program or process that was blocked by the semaphore.

Finally we want to test that a program can close a semaphore if no process is blocked by it and because it is possible to close a semaphore and make it available again we want to test that obtaining every semaphore possible by creating 128 semaphores which is our limit by implementation, closing 128 should then free 128 semaphores again.

`'semaphore_test.c'` is the initprog in our testcase and is responsible for creating new process that have their own thread - This is needed to test if a semaphore is blocked correctly because a process cannot block and unblock itself since that would defeat the purpose.

`'semaphore_test.c'` starts by creating a semaphore it then creates 15 new processes(`'semaphore_test2.c'`) that will use the semaphore created and block themselves.

To test that these are blocked a third program is started(`'semaphore_test3.c'`). This will try to close the semaphore that `'semaphore_test2.c'` is blocked by, the test is, that it should not be allowed to close the semaphore before every process(`'semaphore_test2.c'`) has left the sleep_queue.

For every time the while loop in `'semaphore_test3.c'` runs, `'semaphore_test3.c'` calls `'userland_semaphore_V'` and decreases the value by one and unblocks one

process('semaphore_test2.c').

Finally it exits and 'semaphore_test.c' continues to the next test.

The next test tries to create 129 semaphores which should result in a failure once it tries to create number 129, the rest should be created with success. Once the 129 has been created with or without failure, we test that they can also be destroyed by destroying the 128 semaphores we have just created, we loop through 129 entries, but the number 129 semaphore does not exist thus it results in an error which is intended.

Finally we try to create 128 semaphores as before to test that all was closed with success, we then close them immediatly after since we stated previously that it was important to close them for the sake of the buenos system.

In case of failure in our test the system will halt.

In our enviroment the tests run through every test with success.